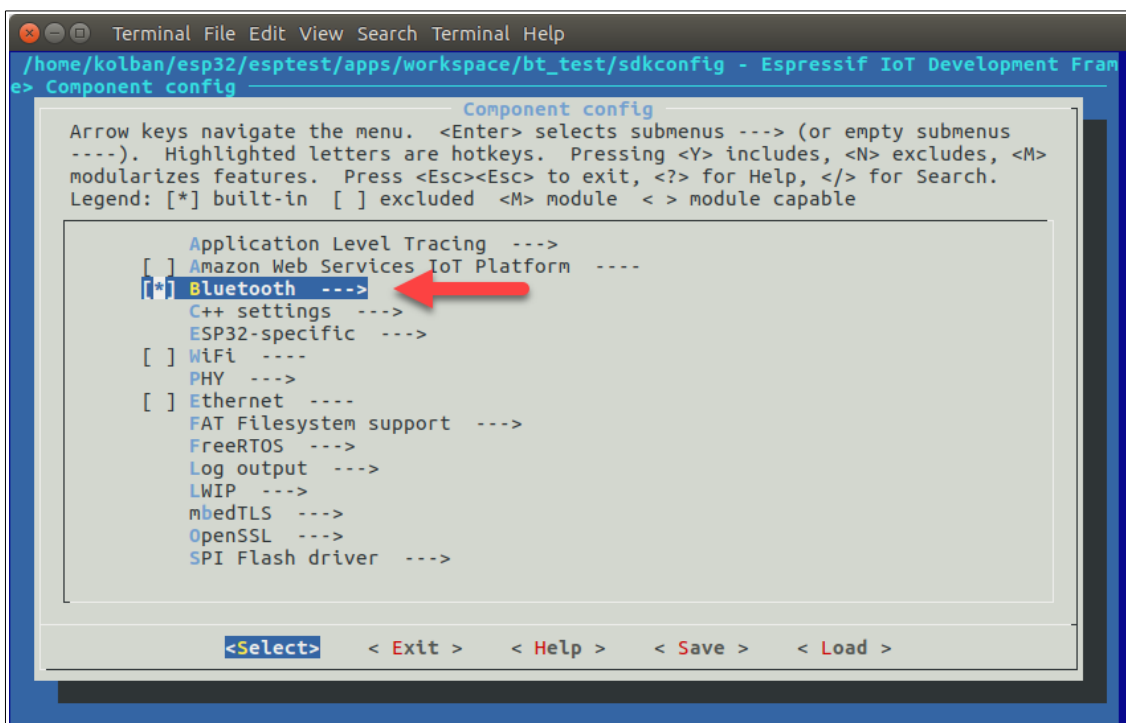


Bluetooth BLE

Using the low level ESP-IDF functions to build a BLE solution is considered difficult. There are a lot of concepts and there is a lot of state to be considered plus events needing to be handled. This feels like an ideal candidate for a higher level class encapsulation to simplify our tasks. To that end, we have created a set of C++ classes that provide all the needed BLE capabilities with, what we hope to be, an easier model of usage.

When using these classes, we must use `make menuconfig` to configure our ESP-IDF environment.

First, we need to ensure that Bluetooth is enabled. Start `make menuconfig` and navigate to Component config and select Bluetooth:



Next drill into the Bluetooth entry and ensure "Bluebird Bluetooth stack enabled" is selected:

```
Terminal File Edit View Search Terminal Help
/home/kolban/esp32/esptest/apps/workspace/bt_test/sdkconfig - Espressif IoT Development Fram
e> Component config > Bluetooth

Bluetooth
Arrow keys navigate the menu. <Enter> selects submenus ---> (or empty submenus
----). Highlighted letters are hotkeys. Pressing <Y> includes, <N> excludes, <M>
modularizes features. Press <Esc><Esc> to exit, <?> for Help, </> for Search.
Legend: [*] built-in [ ] excluded <M> module <> module capable

--- Bluetooth
[*] Bluebird Bluetooth stack enabled --->

<Select> < Exit > < Help > < Save > < Load >
```

Finally, drill into the "Bluebird Bluetooth stack enabled" and set the "Bluetooth event (callback to application) task stack size" to be 8000 (or more).

```
Terminal File Edit View Search Terminal Help
/home/kolban/esp32/esptest/apps/workspace/bt_test/sdkconfig - Espressif IoT Development Fram
e> Component config > Bluetooth > Bluebird Bluetooth stack enabled

Bluebird Bluetooth stack enabled
Arrow keys navigate the menu. <Enter> selects submenus ---> (or empty submenus
----). Highlighted letters are hotkeys. Pressing <Y> includes, <N> excludes, <M>
modularizes features. Press <Esc><Esc> to exit, <?> for Help, </> for Search.
Legend: [*] built-in [ ] excluded <M> module <> module capable

--- Bluebird Bluetooth stack enabled
(8000) Bluetooth event (callback to application) task stack size
[ ] Bluebird memory debug
[ ] Classic Bluetooth
[ ] Release DRAM from Classic BT controller
[*] Include GATT server module(GATTS)
[*] Include GATT client module(GATTC)
[*] Include BLE security module(SMP)
[ ] Close the bluebird bt stack log print
(4) BT/BLE MAX ACL CONNECTIONS(1~7)

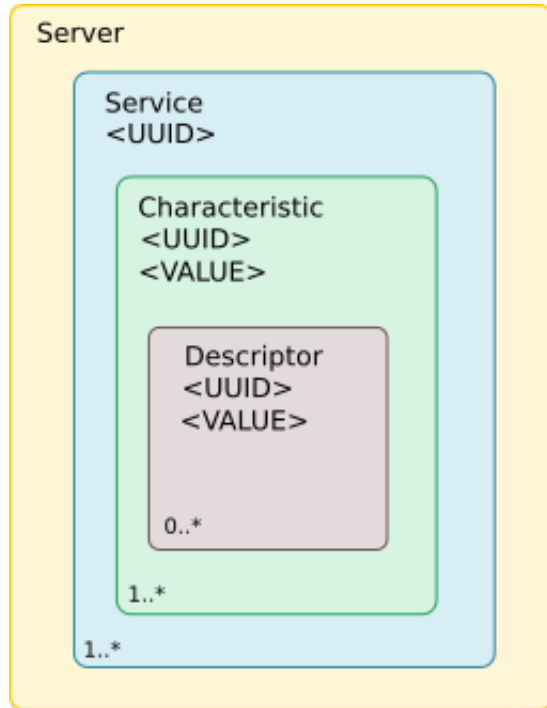
<Select> < Exit > < Help > < Save > < Load >
```

Save all your changes are rebuilt the source.

We'll split our story into two parts ... one being a BLE Server and the other being a BLE Client.

A BLE Server

A BLE Server is also known as a BLE Peripheral. It is likely that this is how you will most commonly use the ESP32. With a BLE Server, you will expose one or more services where each service has one or more characteristics and each characteristic may have zero or more descriptors. I think we'll agree that is quite a lot going on.



In our C++ model, we create the concept of classes that represent these items.

- `BLEServer` – Models a server.
- `BLEService` – Models a service. Owned by a `BLEServer`.
- `BLECharacteristic` – Models a characteristic. Owned by a `BLEService`.
- `BLEDescriptor` – Models a descriptor. Owned by a `BLECharacteristic`.

And also:

- `BLEAdvertising` – Models advertising. Owned by a `BLEServer` to let others know of our existence.

At a high level, the pseudo code of a minimal BLE server becomes:

```
// Initialize the BLE environment
BLE::initServer("ServerName");

// Create the server
BLEServer* pServer = new BLEServer();

// Create the service
BLEService* pService = pServer->createService(ServiceUUID);

// Create the characteristic
BLECharacteristic* pCharacteristic =
    pService->createCharacteristic(CharacteristicUUID, properties);
// Set the characteristic value
pCharacteristic->setValue("Hello world");

// Start the service
pService->start();
```

Hopefully you see how this fits together ... we create the server, we create the service, create a characteristic upon the service, set a value for the characteristic and then ask the service to start responding to requests.

If you are familiar with the BLE APIs, you may notice what is missing ... all the complexity and glue code necessary for event handling and processing of BLE requests. All of this is *handled for you* by the implementation of the classes. This means that you can focus on your intent/usage of BLE while keeping your *plumbing* to a minimum.

The goal of these classes is to efficiently process BLE workflow while encapsulating the plumbing so you don't need to worry about it ... but within is the danger that the implementation will restrict you from some tasks by hiding functions that you might otherwise have needed for your own project. Thankfully, this has not been seen to be the case. The classes expose simple high level APIs for the 95% of common practices while at the same time providing methods that can be called to tweak and tailor the operations for the rarer cases. Hopefully you won't need those often but, if and when you do, they are there for you.

When a BLE Server is running, what must happen next is that peer devices (clients) must be able to locate it. This is made possible through the notion of advertising. The BLE Server can broadcast its existence along with sufficient information to allow a client to know what services it can provide.

Once we have started the BLE Server, we can ask it for an object (`BLEAdvertising`) that owns the advertisements that the server produces:

```
BLEAdvertising* pAdvertising = pServer->getAdvertising();
pAdvertising->start();
```

Once performed, the server can start to be dynamically found by the clients. Of course a server doesn't *need* to advertise. If a client should otherwise be informed (or remember) the address of the BLE server, it can request a connection at any time.

A core concept of BLE is the notion of the characteristic. Think of this as a stateful *record* that has an identity and a value. A peer device (if permitted) can read the value of the characteristic or set a new

value. Remember, it is the BLE characteristic that owns the existence of the value so that it may be served up or changed upon request. This is the core notion of BLE. The value read from the characteristic provides information. For example, if the characteristic represents your heart-rate measured from a sensor, then a remote client can retrieve your current heart rate by reading the current characteristic value. In this case, a remote client will not be setting the value, instead the value will be changed internally by the server either each time a read request by a client is made or when a new sensor reading is taken. Alternatively, the characteristic maintained in the BLE Server may represent the state of something that can be activated or changed. As another example, imagine that the characteristic represents the state of your car door's lock. You can read the characteristic to determine that your door is locked when you leave or, conversely, you can set the characteristic's value to an unlocked state which would result in the server mechanically unlocking the door.

To model these notions in our C++ classes, we leverage the notion that a C++ class can be sub-classed to be more specialized. A class called `BLECharacteristicCallbacks` provides two methods that can be over-riden:

- `onRead(BLECharacteristic* pCharacteristic)` – Called when a read request arrives from a client. A new value for the characteristic can be set before return from the function and will be used as the value received by the client.
- `onWrite(BLECharacteristic* pCharacteristic)` – Called when a client initiated write request arrives. The new value has been set in the characteristic already.

To utilize, we can then override these functions in our own C++ class that subclasses the one provided:

```
class MyCallback: public BLECharacteristicCallbacks {
    void onRead(BLECharacteristic*) {
        // Do something before the read completes.
    }

    void onWrite(BLECharacteristic*) {
        // Do something because a new value was written.
    }
}
```

To use this technique, we inform our `BLECharacteristic` about a callback handler. For example:

```
pCharacteristic->setCallbacks(new MyCallback());
```

Here is an example that sends the time since startup each time a client requests the value:

```
class MyCallbackHandler: public BLECharacteristicCallbacks {
    void onRead(BLECharacteristic* pCharacteristic) {
        struct timeval tv;
        gettimeofday(&tv, nullptr);
        std::ostringstream os;
        os << "Time: " << tv.tv_sec;
        pCharacteristic->setValue(os.str());
    }
}
```

```
}  
};
```

Similar to the characteristic callbacks, we also have server callbacks. These inform about client connection and disconnection events. These are subclassed from the `BLEServerCallbacks` class which has virtual methods for:

- `onConnect(BLEServer* pServer)` – Called when a connection occurs.
- `onDisconnect(BLEServer* pServer)` – Called when a disconnection occurs.

These could be used to enable or disable sensor readings. For example, if we have no clients connected then there is no need to spend energy sampling a value if there is no-one there to read it. However, when a client connects, we can detect that and start reading from the sensor at that point until a subsequent disconnection indication is detected.

Another consideration for our BLE Server is the idea that we might want to "push" data to the peer when something interesting happens. Up until now we have considered the idea that the server can receive read requests to get the current value or can receive write requests to set the current value. There is one more operation that we are interested in that is invoked on the server side by the server and asynchronously to the client. That operation is called "indicate". It is used to signal (or indicate) to the client that the characteristic's value has changed. The client will receive an indication event to let it know that the change has occurred.

On the server, we can cause an indication to occur by invoking:

```
pCharacteristic->indicate();
```

The current value of the characteristic will be the value transmitted to the peer. We might pair this call with a previous request to set a new value:

```
pCharacteristic->setValue("HighTemp");  
pCharacteristic->indicate();
```

A similar function to `indicate()` is called `notify()`. The distinction between them is that an `indicate()` receives a confirmation while a `notify()` does not receive a confirmation.

Associated with the idea of indications/notifications, is the architected BLE Descriptor called "[Client Characteristic Configuration](#)" which has UUID 0x2902. This contains two distinct bit fields that can be on or off. One bit field governs Notifications while the other governs Indications. If the corresponding bit is on, then the server can/may send the corresponding push. For example if the Notifications bit is on, then the server can/may send notifications. The primary purpose of the descriptor is to allow a partner to request that the server *actually* send notifications or indications. Here is an example. Imagine that we have a BLE Server that can publish notifications when data changes. Now imagine that a BLE Client connects but is actually not interested in receiving these. If the BLE

Server executes a notification push and sends the data, that will be wasted effort/energy as the client doesn't want or can't use the information. What we really want is that the client should inform the server that *if* it wants to push new data then it either may or shouldn't. And that is where this descriptor comes into play. If the descriptor is present on a characteristic then then client can remotely change the bit flag to enable or disable notifications and indications. Since the descriptor flag is stored local to the server, the server is at liberty to examine the flag before performing a radio transmission. What this means is that the client/peer can toggle on or off its desire to receive notifications and the server should honor those requests.

The BLE specification constrains the maximum amount of data that can be sent through a notification or indication to be 20 bytes or less. Take this into account in your designs. If the value of a characteristic is greater than this amount, then only the first 20 bytes of the data will be transmitted.

It is likely that your own BLE server application is going to expose its own set of characteristics. Through a characteristic you can set and get the value as a binary piece of storage, however this may be too low level for you. An alternative is to utilize the features of the C++ language and model your own specialized characteristic as a sub class of `BLECharacteristic`. This could then encapsulate the lower level value's getter and setter with your own customized version.

For example, if your characteristic represented a temperature, you could create:

```
class MyTemperatureCharacteristic: public BLECharacteristic {
    MyTemperature(): BleCharacteristic(BLEUUID(MYUUID)) {
        setTemperature(0.0);
    }

    void setTemperature(double temp) {
        setValue(&temp, sizeof temp);
    }

    double getTemperature() {
        return *(double*)getValue();
    }
}
```

A BLE Client

Now we will turn our attention to the second half of our story, namely that of being a client. In this story, our ESP32 doesn't host services but instead wishes to be a consumer of services hosted elsewhere. The story can further be broken down into two sub-stories ... namely scanning and interaction.

If we assume that our ESP32 starts up and wishes to be a client of a remote BLE server then it has to connect to that server. In order to connect to the server, we need to know the address of the target server. An address is a 6 byte value commonly written in the form

```
nn:nn:nn:nn:nn:nn
```

While in principle this can be hard-coded into your application or manually entered through some interaction story, this is not the common practice. Instead, we perform a procedure known as a *scan*.

Scanning is the idea that we actively listen on the BLE radio frequencies for servers which are actively advertising their existence. As an analogy, imagine we enter a dark room and we have no idea who is in there with us. We take off our ear-muffs (yes, for some reason we entered a dark room with no light at all wearing ear-muffs ... but then, this is just an analogy) and we start hearing voices. We hear "Hi this is Bob, I can make toast" and "Hi this is Susan, I can tell you what temperature it is" and we also hear "Hi this is Brian, ask me what I can do". In this analogy, Bob, Susan and Brian are the addresses of the BLE devices. Each of them are continually saying out loud that they exist and, in some cases, what they can do for us.

When we perform a BLE scan, we receive short size records of data that always contain the address of the advertiser and *sometimes* additional information such as the services they provide or other descriptive items. This information arrives at us passively. All we need do is listen and we learn. Should we need to learn more about the devices, we can connect to them (even the ones that told us little) and explicitly ask them for more details of what they can do.

And that is the principle of scanning. In our C++ BLE story, we have modeled this through the C++ class called `BLEScan`. We get an instance of this class by asking the BLE device for it using:

```
BLEScan* pMyScan = BLE::getScan();
```

The object returned to us is a singleton. This was chosen because we will only want to be scanning once per ESP32 at any given time.

The `BLEScan` object, when presented to us, isn't actually yet scanning, to start it scanning we invoke its `start()` method passing in the duration of how long we would like it to scan for. This is an interval measured in seconds. For example:

```
pMyScan.start(30); // Scan for 30 seconds
```

This is a blocking call. It will return after the scan period has elapsed. It returns an object instance of type `BLEScanResults` that contains the information for each unique device found. If we have an instance of `BLEScanResults` we can ask it how many results it contains (via a call to `getCount()`) and we can then retrieve each advertised device (via a call to `getDevice(index)`).

Since we are scanning, we will want to know as soon as possible about the devices that we find and this is where a callback function comes into play. Before calling `start()` we can register a callback function that will be invoked for each peer device that was found. An abstract C++ class called `BLEAdvertisedDeviceCallbacks` works for us here. This has a method on it called `onResult()` which will be invoked for each unique result found during scanning. Note that we don't pass the same detected device twice; if it wasn't the one you were previously looking for, it still won't be the one you want a short while later. This method is passed in an instance to an object of type `BLEAdvertisedDevice` that describes the nature of the device that we found.

For example:

```
class MyCallbacks: public BLEAdvertisedDeviceCallbacks {
    void onResult(BLEAdvertisedDevice advertisedDevice) {
        // Do something with the found device ...
    }
}

BLE::initClient();
BLEScan* pMyScan = BLE::getScan();
pMyScan->setAdvertisedDeviceCallbacks(new MyCallbacks());
pMyScan->start();
```

What this means is that for every unique device that the scan finds, a call to `onResult()` will be made which is passed an instance of a `BLEAdvertisedDevice` which describes that device. Now let us look at what a device *may* tell us. The only thing that a device will tell us for sure is its BLE address. Beyond that, everything else is optional. The *possible* attributes that an advertised device may inform us about is large and the following subset have been modeled (so far):

- Appearance
- Manufacturer data
- Name
- RSSI
- Primary service UUID
- Transmit power

Since none of these are mandatory in an advertisement, for any given advertised device we learn about, we can't immediately ask about the value of the property. Instead, we have methods that tell us which properties are present and, if present, **then** we can ask for its value.

For example, in our processing logic we may code:

```
if (advertisedDevice.hasServiceUUID()) {
    BLEUUID service = advertisedDevice.getServiceUUID();
    if (service.equals(BLEUUID((uint16_t)0x1802)) {
        // we found a useful device ...
    }
}
```

(The above is merely an example of the logic that can be employed).

If we wish to end the scan early; presumably because we found the device we wanted, we can call the `stop()` method of the `BLEScan` object. A reference to the `BLEScan` is conveniently available within the `BLEAdvertisedDevice` object.

```
advertisedDevice.getScan()->stop();
```

From all of this ... we end up with the identification of a device that was of interest to us and since we want to be a BLE Client, this will presumably mean the device to which we wish to connect. The key item in the advertisement now becomes the device's address which we can obtain through a call to `getAddress()`.

Now we can turn our attention to actually connecting to the server.

A BLE Client is modeled as the `BLEClient` class. We obtain an unconnected instance of this by asking the ESP32 for one:

```
BLEClient* pMyClient = BLE::createClient();
```

Next we request a connection to the target device.

```
pMyClient->connect(address);
```

... and that's it. The `connect()` is a blocking call and, on return, we are connected. However, that isn't the end of the story. Just connecting to a peer is usually not enough. Now we want to read and write the values that the remote BLE server is hosting. If we think back to our understanding of the story, the remote server has one or more services and each service has one or more characteristics and it is these characteristics that hold the values. This means that it doesn't make sense to say "I want to read the value of the remote BLE server" ... instead we must say "I want to read the value of a named remote characteristic". However, even that is not enough. There may be multiple services on the BLE server each of which have their own characteristic which may be of the same characteristic type and hence not be uniquely identified. Thus we end up with the final concept of "I want to read the value of a named remote characteristic that is owned by a named service". From this notion, we now get to introduce two more models. Those are the `BLERemoteService` and `BLERemoteCharacteristic`.

Once we have connected to a BLE Server, we can request a reference to the `BLERemoteService` that models a service on that server:

```
BLERemoteService* pMyRemoteService = pClient->getService(serviceUUID);
```

and once we have a reference to the service, we can ask *that* service for a reference to the desired characteristic.

```
BLERemoteCharacteristic* pMyRemoteCharacteristic =  
    pMyRemoteService->getCharacteristic(characteristicUUID);
```

and ... finally ... we can work with the values.

```
std::string myValue = pMyRemoteCharacteristic->readValue();
```

to read the value and ...

```
pMyRemoteCharacteristic->writeValue("abc");
```

to write a new value. If all that seems like a lot of work ... lets put it together in context and see:

```
// Create the client  
BLEClient* pMyClient = BLE::createClient();
```

```
// Connect the client to the server  
pMyClient->connect(address);
```

```
// Get a reference to a specific remote service on the server
BLERemoteService* pMyRemoteService = pClient->getService(serviceUUID);

// Get a reference to a specific remote characteristic owned by the service
BLERemoteCharacteristic* pMyRemoteCharacteristic =
    pMyRemoteService->getCharacteristic(characteristicUUID);

// Retrieve the current value of the remote characteristic.
std::string myValue = pMyRemoteCharacteristic->readValue();
```

... and that's it. Hopefully you see that is pretty elegant (and if not, *please* contact me and we'll see if we can't improve on it ... but come bearing suggestions for what can be improved).